

# Power BI refreshes reinvented

Streamlining Power BI Refreshes Advanced  
Methods for Real-Time Refreshes

# Thank you, partners



# After this session

## Storage modes

Understand the different storage modes in Power BI and how this relates to refreshes.

## Refresh operations

Understand which different options there are to refresh your Semantic Model manually as well as in an automated fashion.

## Partitioning

Understand how different partitions in the data model are generated and can be triggered to refresh.

## Orchestration

Understand how you can orchestrate your Power BI refresh and bind to other processes.

# Marc Lelijveld

Technical Evangelist | Solution Architect  
Macaw Netherlands



@MarcLelijveld



linkedin.com/in/MarcLelijveld



Data-Marc.com



DutchFabricUsergroup.com

FAVORITE STUFF:



# What we cover today

- Refreshes?
- Incremental refresh
- High frequency refreshes & hybrid tables
- Refresh individual objects
- Orchestration & automation
- Wrap-up





# REFRESH




## Refreshes




# Understanding storage modes



Import



Dual



DirectQuery

Oh wait... we now also have DirectLake!

Storage Modes - Power BI Desktop

Search

Marc Lelijveld

File Home Help External Tools

Clipboard Data Queries Relationships Security Q&A Language Linguistic schema Sensitivity Share

Product Subcategory

- English Product Subcategory N...
- French Product Subcategory Na...
- ProductCategoryKey
- ProductSubcategoryAlt...
- ProductSubcategoryKey
- Spanish Product Subcategory N...

Collapse ^

Product

- Arabic Description
- Chinese Description
- Class
- Color
- Days To Manufacture
- Dealer Price

Collapse ^

Internet Sales - Agg

- Count
- Order Calendar Year
- ProductSubcategory/Key
- Sales Amount

Collapse ^

Internet Sales

- Carrier Tracking Number
- CurrencyKey
- Customer PO Number
- CustomerKey
- Discount Amount
- DueDateKey
- Extended Amount
- Freight Amount
- Order Calendar Year

Collapse ^

Properties

General

Name

Internet Sales

Description

Enter a description

Synonyms

Enter a comma-separated list of synonyms for Q&A

Row label

Select a row label

Key column

Select a column with unique values

Is hidden

No

Is featured table

No

Advanced

Storage mode

DirectQuery

Fields

All tables Layout 1

# Understanding Storage Modes

## Three storage modes

- Import – data cached in the model
- DirectQuery – queries are submitted to the back-end data source
- Dual – can act in both above storage modes, depending on query context

## Configuring storage modes

- Storage modes are set on table level
- Setting storage mode to Import is an irreversible operation
- Data in DirectQuery mode cannot be displayed in the data tab



# Caches and DirectQuery



## Risks of mixing storage modes

- Avoid mismatch in results when query bits DirectQuery compared to Import
- Data cached (import) could be behind compared to DirectQuery data
- Make sure cached data is kept in sync – regularly refresh!

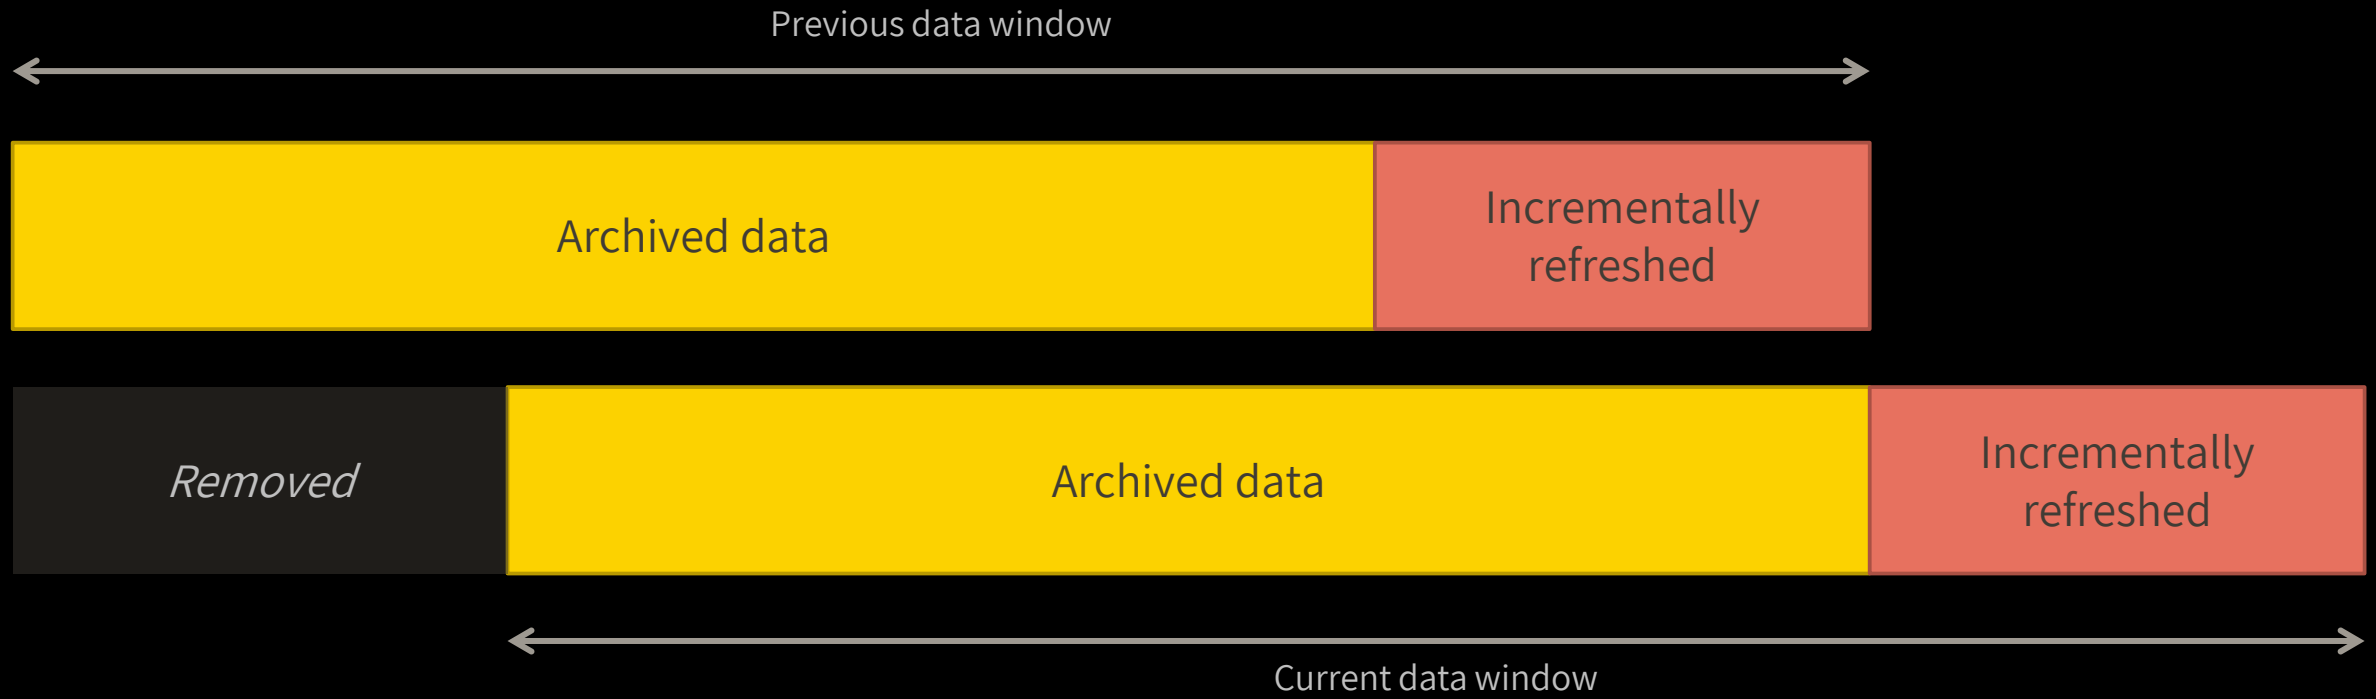
# Benefits to choose one or the other

- Improve query performance – cache data for faster end-user performance
- Data refresh optimization – no need to refresh for non-cached data
- Near-real time requirements – reduce query latency when in DirectQuery mode
- Large Semantic Models – choose to not import certain data by leveraging aggregations for example

# Different refresh operation types

Type	What	Power BI Desktop 	Power BI Service (Interface) 
Manual	Individual Table	✓	✗
	Full model	✓	✓
Scheduled	Individual Table	✗	✗
	Full model	✗	✓






# Incremental refresh

# Considering incremental refresh

## Why should you use it?

- Large data volumes
- Faster refreshes
- Resource consumption reduced  
(Both source & Power BI side)



# Incremental refresh

Incremental refresh is supported for Power BI Premium, Premium per user, **Power BI Pro**, and Power BI Embedded Semantic Models.

Getting the latest data in **real time** with DirectQuery is **only supported for Power BI Premium\***, Premium per User, and Power BI Embedded Semantic Models.


*\* Premium capacities = Fabric capacities*



# Supported data sources

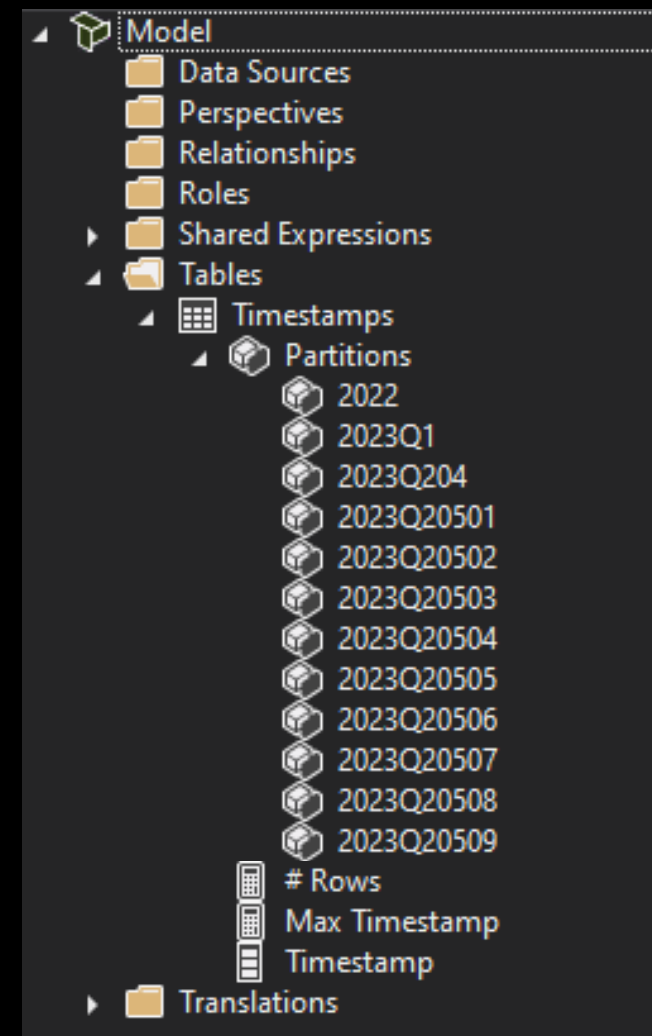
## Data sources must be

- Structured
- Relational
- Allow query folding
- Typically – SQL like sources



# Breaking up your data in partitions

- Each partition represents a time frame
- New partitions get added automatically
- Old partitions (in archive) get merged



# Configure Incremental Refresh

### Manage Parameters

RangeStart

RangeEnd

Name

RangeStart

Description

☒ Required

Type

Date/Time

Suggested Values

Any value

Current Value

1/1/2022 12:00:00 AM

OK

Cancel

### Incremental refresh and real-time data

Refresh large tables faster with incremental refresh. Plus, get the latest data in real time with DirectQuery (Premium only). [Learn more](#)

① These settings will apply when you publish the dataset to the Power BI service. Once you do that, you won't be able to download it back to Power BI Desktop. [Learn more](#)

- Select table**

Orders
- Set import and refresh ranges**

☒ Incrementally refresh this table

Archive data starting 5 Years before refresh date

Data imported from 1/1/2017 to 11/17/2022 (inclusive)

Incrementally refresh data starting 3 Days before refresh date

Data will be incrementally refreshed from 11/18/2022 to 11/20/2022 (inclusive)
- Choose optional settings**

☐ Get the latest data in real time with DirectQuery (Premium only) [Learn more](#)

Selected table cannot be folded for DirectQuery.

☐ Only refresh complete days [Learn more](#)

☐ Detect data changes [Learn more](#)
- Review and apply**

5 years before

Archived

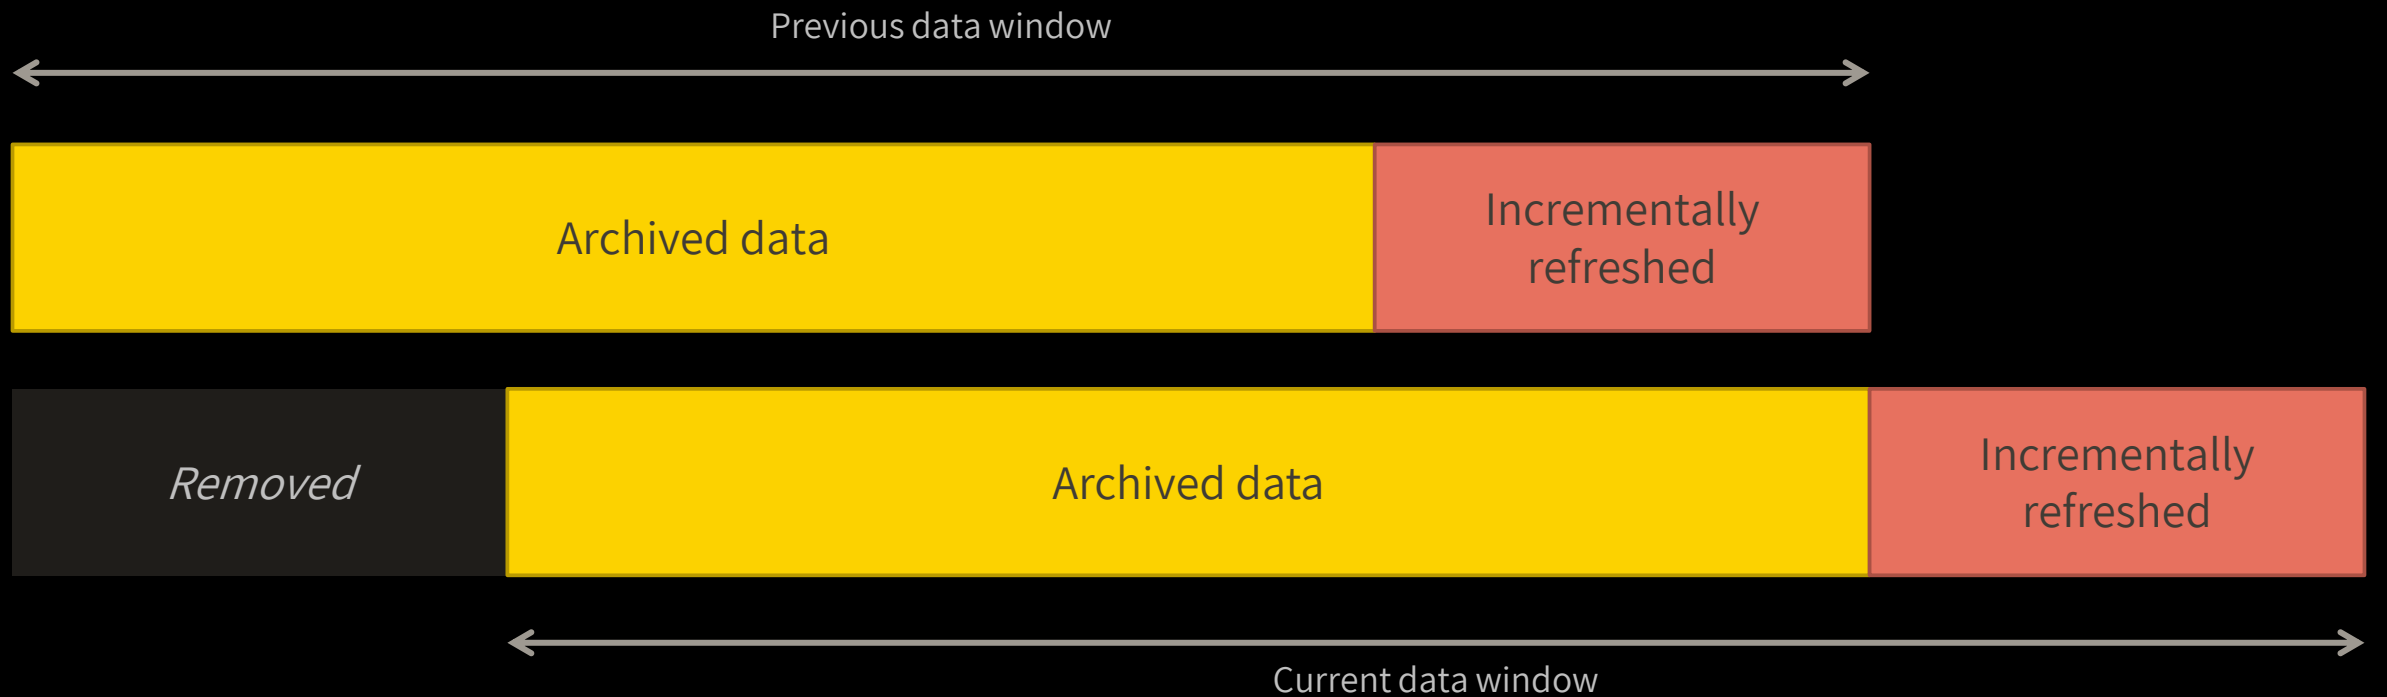
3 days before

Incremental Refresh

Refresh date




# Moving time frame




# Detect data changes

- Even more selective refresh – only that what has been changed
- Based date/time column used to identify and refresh only those days where the data has changed
- Typically, system generated dates like Modified Date Time
- Should never be the same column as your incremental periods
- The maximum value is evaluated for each partition **in the incremental range** to detect whether the partition should be refreshed or not



**Demo** - configure  
incremental refresh





**High frequency refreshes & hybrid tables**

# What is high frequency?

## Let's say real time?

- Hourly?
- Minutes?
- Seconds?

Or...

- Just daily batches



# Refresh limitations

Frequency	Manual	Scheduled (via the service)
Free	Up to 8 times per day	
Pro	Up to 8 times per day	
Premium per User	Unlimited	48 times per day
Premium / Fabric Capacity	Unlimited	48 times per day

# Hybrid tables

- Live / Realtime data in Power BI
- Combines different storage modes on partition level in a single table
- Goes hand-in-hand with Incremental Refresh

Granularity	Name	Row Count	
Year	2011	295,489,717	Archived: <b>Import</b>
Year	2012	297,678,498	
Year	2013	295,575,442	
Year	2014	292,477,875	
Year	2015	297,780,469	
Year	2016	294,060,081	
Year	2017	300,419,682	
Year	2018	296,541,108	
Year	2019	292,787,420	
Year	2020	299,273,979	
Quarter	2021Q1	74,135,277	Incremental refresh: <b>Import</b>
Month	2021Q104	24,939,498	
Day	2021Q10501	820,805	
Day	2021Q10502	826,885	Real time: <b>DirectQuery</b>
Day	2021Q10503	821,043	
Day-DirectQuery	2021Q10504-DQ	271,110	
Total		3,063,898,887	

# Hybrid tables

- Implementation with Incremental Refresh
- Customizable via 3rd party tooling like Tabular Editor (*unsupported*)

## Limitation:

Only 1 DQ partition per table allowed.

Incremental refresh and real-time data

ⓘ

These settings will apply when you publish the dataset to the Power BI service. Once you do that, you won't be able to download it back to Power BI Desktop. [Learn more](#)

1. Select table

Sales Agg

2. Set import and refresh ranges

☒

Incrementally refresh this table

Archive data starting 10 Years before refresh date

Data imported from 3/2/2012 to 2/20/2022.

Incrementally refresh data starting 10 Days before refresh date

Data will be incrementally refreshed from 2/20/2022 to 3/2/2022.

3. Choose optional settings

☒

Get the latest data in real time with DirectQuery (Premium only) [Learn more](#)

☒

Only refresh complete days [Learn more](#)

☐

Detect data changes [Learn more](#)

4. Review and apply

Archived

Incremental Refresh

Real time

10 years before refresh date

10 days before refresh date

Refresh date

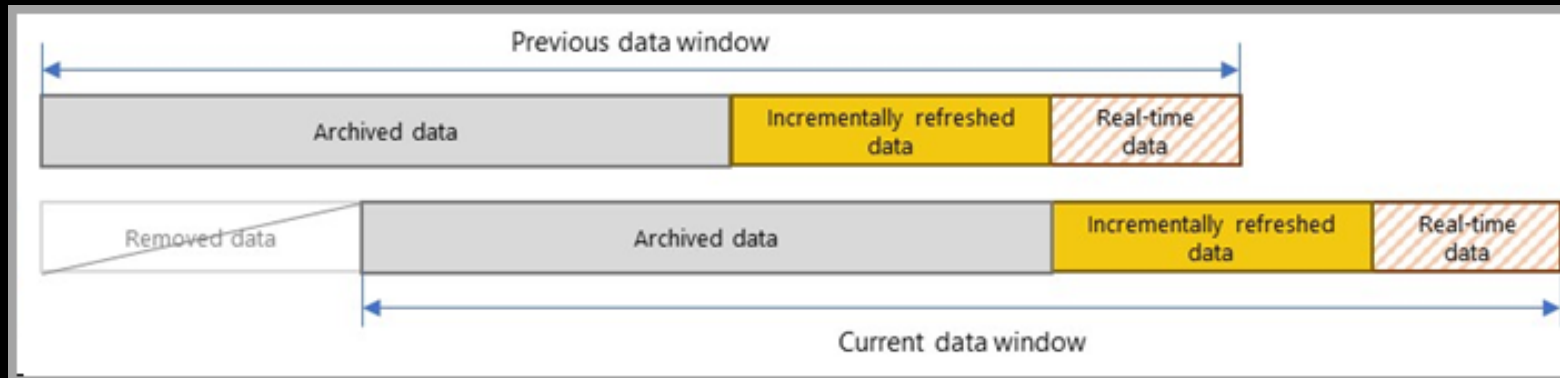
Apply

Cancel



# Hybrid tables – what challenge does it solve?

- Realtime scenarios without full tables on DQ mode
- No complex refresh mechanisms needed with partition refresh and queries over XMLA
- No more multiple tables and complex DAX to combine to achieve the same goal



# Hybrid tables – Keep in mind that...

- Premium feature
- DAX restrictions for DirectQuery apply
- Limited Power Query & DAX capabilities (due to DQ)
- Requires Large Semantic Model Format (storage) in workspace
- Performance hit on upstream data sources



**Demo** - configure  
hybrid tables and  
streaming data

Hybrid Table
...
Feb 2022
March 2022
...
...
...



Import



DirectQuery

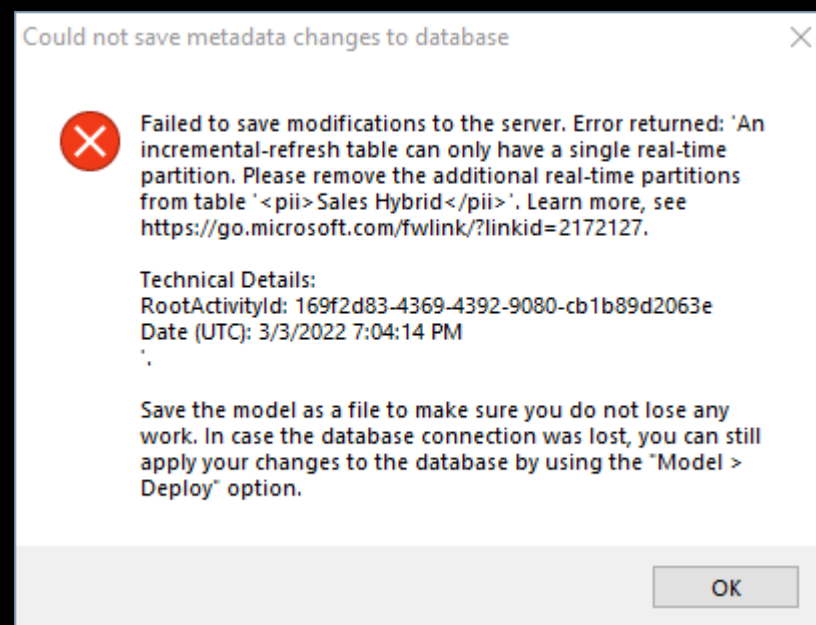


▼	📊	Sales Hybrid
▼	📦	Partitions
	📦	2021Q309
	📦	2021Q410
	📦	2021Q411
	📦	2021Q412
	📦	2022Q101
	📦	2022Q102
	📦	2022Q10301
	📦	2022Q10302
	📦	2022Q10303-onward

# Can I change partition storage modes?

You cannot adjust tables with incremental refresh policies applied. However, there are other options to consider.

▼ <b>Options</b>	
Data View	Default
Mode	DirectQuery
Source Type	PolicyRange





Hybrid Table
...
< 2022
>= 2022
...
...
...

} DirectQuery

} Import

### DirectQuery Partition definition

```
let
    Source = Sql.Databases((SynapseEndpoint)),
    #"WideWorldImportersDW-Standard" = Source{[Name=(DatabaseName)]}[Data],
    Fact_Sale = #"WideWorldImportersDW-Standard"[[Schema=(SynapseSchema),Item="Sales"]][Data],
    #"Filtered Rows" = Table.SelectRows(#"Fact_Sale", each
        [InvoiceDateKey] < #date(2022, 01, 01) and
        [InvoiceDateKey] >= #date(2013, 01, 01)),
        SplitByCharacter = Table.TransformColumnNames(#"Filtered Rows", each
            Text.Combine(
                Splitter.SplitTextByCharacterTransition({"a".."z"}, {"A".."Z"})(_)
                , " ")
        )
in
    SplitByCharacter
```

### Import Partition definition

```
let
    Source = Sql.Databases((SynapseEndpoint)),
    #"WideWorldImportersDW-Standard" = Source{[Name=(DatabaseName)]}[Data],
    Fact_Sale = #"WideWorldImportersDW-Standard"[[Schema=(SynapseSchema),Item="Sales"]][Data],
    #"Filtered Rows" = Table.SelectRows(#"Fact_Sale", each
        [InvoiceDateKey] >= #date(2022, 01, 01) and
        [InvoiceDateKey] < #date(2022, 03, 31)),
        SplitByCharacter = Table.TransformColumnNames(#"Filtered Rows", each
            Text.Combine(
                Splitter.SplitTextByCharacterTransition({"a".."z"}, {"A".."Z"})(_)
                , " ")
        )
in
    SplitByCharacter
```



**Refresh individual objects**

# Effective refreshing

## Considerations

- Refreshing the entire model takes too long with high load on sources
- Can we only refresh certain tables?
- Can we only refresh certain partitions?
- Can we use DQ tables/partitions (Hybrid Tables)

## What do we need?

- Enhanced refresh API
- XMLA Endpoints

# Smaller increments



Bind your refresh operations in Power BI to refresh operations in for example your data platform



Run smaller increments to speed up the end-to-end process



Trigger one (or multiple) selective table(s) or partition(s) at a time.

# Smaller increments

Smaller increments can only be triggered to refresh via:

- Manual processing via Management Studio (or other tools)
- Power BI Enhanced Refresh API - Previously known as Async refresh API
- XMLA endpoints

# Enhanced refresh API

## More granular controls

- type - full / clearValues / calculate / dataonly / automatic / defragment
- Commitmode – default is transactional
- maxParallism – 10 max
- retryCount – number of retries before failing
- objects – which objects to refresh
- applyRefreshPolicy – true/false whether you want to apply potential incremental refresh if configured
- effectiveDate – in case of incremental refresh, this parameter overwrites the current date



# Enhanced refresh API

Specify the objects to refresh

## POST

<https://api.powerbi.com/v1.0/myorg/groups/f089354e-8366-4e18-aea3-4cb4a3a50b48/datasets/cfafbeeb1-8037-4d0c-896e-a46fb27ff229/refreshes>

```
{
  "type": "Full",
  "commitMode": "transactional",
  "maxParallelism": 2,
  "retryCount": 2,
  "objects": [
    {
      "table": "DimCustomer",
      "partition": "DimCustomer"
    },
    {
      "table": "DimDate"
    }
  ]
}
```



**Demo** – Refresh one table or partition

# More than refreshes

Enhanced refresh API also allows you to

- **GET** refresh history
- **GET** individual refresh details like progress or failures
- **DELETE** a running refresh

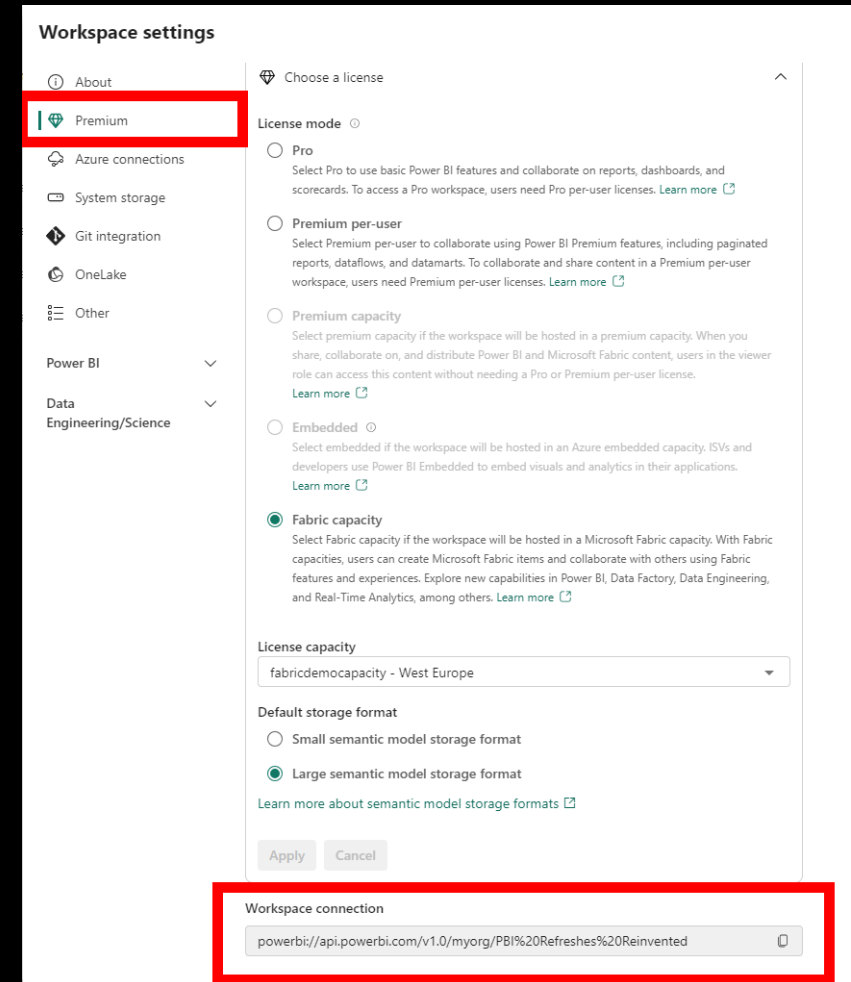

JSON


Copy

```
{
  "startTime": "2020-12-07T02:06:57.1838734Z",
  "endTime": "2020-12-07T02:07:00.4929675Z",
  "type": "Full",
  "status": "InProgress",
  "currentRefreshType": "Full",
  "objects": [
    {
      "table": "DimCustomer",
      "partition": "DimCustomer",
      "status": "InProgress"
    },
    {
      "table": "DimDate",
      "partition": "DimDate",
      "status": "InProgress"
    }
  ]
}
```

# Is the API the only option?

- Execute operations over XMLA endpoints
- Using Tabular Model Scripting Language
- More granular control using SQL Service Management Studio






# Orchestration & Automation

Binding refreshes to processes

# Central E2E orchestration

- Combine pipelines from Data Platform with Power BI
- Lowest latency between source and report
- Incremental loading where possible
- Consider including backup operations for Power BI

# What to use to trigger and bind together?






# What is Semantic Link exactly?

Semantic Link is a feature in Microsoft Fabric that allows you to connect from Synapse Data Science Notebooks to Power BI Semantic Models.

This feature **only** exists and works in Microsoft Fabric.



# Semantic Link in Fabric notebooks!

Refresh your semantic models via a notebook and trigger dependent actions.

- Given Semantic Link uses the REST APIs, you can orchestrate not only your semantic model refresh, but also trigger upstream dataflows for example
- Refresh individual tables, partitions or reprocess partitions through enhanced refresh API
- Anything else what is possible with the REST API



Python

Copy

```
PowerBIRestClient(token_provider: TokenProvider | None = None)
```



## **Demo** – Fabric Notebooks for automation

# Wrap up

LET'S  
RECAP...

Refreshes can be manual or **automated**

Incremental refresh makes your refresh **more efficient**

With **detect data changes**, further optimization can be done

Partitioning in incremental refresh should be **based on functional dates**

**Lower latency** by linking Power BI refreshes to other processes

Use the enhanced refresh API to **trigger individual model objects**

Use the toolkit of your preference for **end-to-end orchestration & automation**

# Resources

## **Setting up scheduled refreshes**

<https://learn.microsoft.com/en-us/power-bi/connect-data/refresh-scheduled-refresh>

## **Configure incremental refresh**

<https://learn.microsoft.com/en-us/power-bi/connect-data/incremental-refresh-configure>

## **Enhanced refresh API**

<https://learn.microsoft.com/en-us/power-bi/connect-data/asynchronous-refresh>

## **Semantic Model connectivity with XMLA endpoints**

<https://learn.microsoft.com/en-us/power-bi/enterprise/service-premium-connect-tools>

## **Tabular Model Scripting Language (TMSL)**

<https://learn.microsoft.com/en-us/analysis-services/tmsl/tabular-model-scripting-language-tmsl-reference?view=asallproducts-allversions>

## **These slides**

<https://github.com/marclelijveld/Slide-decks>


# Thanks for attending!



**Marc Lelijveld**

Technical Evangelist | Architect  
Macaw Netherlands



 @MarcLelijveld

 linkedin.com/in/MarcLelijveld

 Data-Marc.com



[https://bit.ly/dMC2024\\_SessionFeedback](https://bit.ly/dMC2024_SessionFeedback)

**Feedback!**